

# Retrieval-Augmented Generation (RAG) Documentation

## Retrieval-Augmented Generation (RAG) Documentation

### Overview of RAG

Retrieval-Augmented Generation (RAG) combines document retrieval and response generation to produce contextually relevant answers. It uses a knowledge base to retrieve relevant documents and a large language model to generate responses.

### Key Components in the Code

1. **Vector Store (Pinecone):** Stores embeddings of the knowledge base documents for efficient similarity searches.
2. **Embeddings (OpenAI):** Converts text into high-dimensional vectors for similarity matching.
3. **Document Chain:** Combines multiple retrieved documents into a coherent response.
4. **Retrieval Chain:** Executes the retrieval and response generation pipeline.

### Workflow

1. **User Query:** The chatbot receives input via WebSocket.
2. **Document Retrieval:** Retrieves the top relevant documents from Pinecone.
3. **Response Generation:** The LLM generates contextually accurate responses.
4. **Response Delivery:** Sends the generated response or profile to the user.

### Features

1. Top-k Document Retrieval for relevance.

# Retrieval-Augmented Generation (RAG) Documentation

2. Context Awareness through chat history.
3. Dynamic Response Generation guided by retrieved data.

## Limitations

1. Dependency on the knowledge base for accuracy.
2. Scalability concerns with in-memory chat history.

## Future Enhancements

1. Incorporate hierarchical retrieval strategies.
2. Use a database for persistent chat history.
3. Explore fine-tuned embeddings for improved accuracy.